

2022083436 주예지

운영체제 Project 2 Wiki

A Simple Kernel-Level Thread

I. Design:

1.) system call clone(), join()

- The clone() system call must create a new thread sharing the parent's address space and resources but with an independent execution context and stack. The modifications ensure threads share one page table (true threading), correctly set up register state, and avoid lock races or lost wakeups.
- The join() system call must wait for any thread of the calling process to terminate, retrieve its user stack pointer, clean up the kernel TCB, and return the thread's PID. Modifications prevent lost wakeups and ensure safe user-space pointer copy.

2.) thread_create() and thread_join()

- The thread_create() must allocate a user stack and invoke the clone() system call.
- I've allocated a page size user stack by calling malloc(), and called clone() directly.
- The thread_join() must block until a thread exists, retrieve and free its user stack, and return the thread's pid.
- I've added a join() call and made it return -1 if pid<0.

II. Implementation

1.) clone()

In syscall.c, I linked the SYS_clone system call to sys_clone().

```

27  #define SYS_fcfsmode 26
28  #define SYS_yield 27
29  #define SYS_clone 28
30  #define SYS_join 29
31
104  extern uint64 sys_getppid(void);
105  extern uint64 sys_clone(void);
106  extern uint64 sys_join(void);
107
138  [SYS_clone] sys_clone,
139  [SYS_join] sys_join,

```

sysproc.c

Fetches the function pointer, arguments, and stack pointer from the user program and calls the clone() function.

The argaddr() function is used to fetch the arguments passed from user space. It reads the address of each argument and stores them in local variables.

fcn: is the function pointer where new thread starts executing

arg1 and arg2: are the arguments to be passed to fcn

stack: is the address of the stack allocated for the new thread, which must be page-aligned and large enough for the new thread's stack

```
uint64
sys_clone(void){
    void (*fcn)(void *, void *);
    void *arg1, *arg2, *stack;

    // Fetch the function pointer, arguments, and stack address from user space
    argaddr(0, (uint64*)&fcn); // Fetch the function pointer
    argaddr(1, (uint64*)&arg1); // Fetch arg1
    argaddr(2, (uint64*)&arg2); // Fetch arg2
    argaddr(3, (uint64*)&stack); // Fetch stack address

    return clone(fcn, arg1, arg2, stack); // Call real clone
}
```

Below is the implementation of int clone() function in proc.c.

This is to allocate a new struct proc.

```
//allocation, new struct proc | You, 1초 전 • Uncommitted cha
struct proc *np;
struct proc *p = myproc();
```

Threads share the same virtual memory. So instead of copying parent's address space, I am going to point to the same pagetable.

```
// Share
np->pagetable = p->pagetable;
// not copying parent's address space
// but pointing to the same pagetable
np->sz = p->sz; // share memory size
```

epc: exception program counter. The thread will be start running.

sp: set to the top of the user stack

a0 and a1 are the first two arguments passed by registers.

This sets up the thread to run fcn(arg1, arg2) with its own stack.

```
// Copy trapframe
*(np->trapframe) = *(p->trapframe); //copy parent's trapframe
np->trapframe->epc = (uint64)fcn; // start address
np->trapframe->sp = (uint64)stack + PGSIZE; // top of user stack
np->trapframe->a0 = (uint64)arg1;
np->trapframe->a1 = (uint64)arg2;
```

Then copy the file descriptors, and the thread's parent is set to the creating thread.
And then I set the state to RUNNABLE, and return the calling thread, the new thread's PID.

```
// copy file descriptor
for (i = 0; i < NOFILE; i++)
    if (p->ofile[i])
        np->ofile[i] = filedup(p->ofile[i]); //filedup: duplicated file descriptor
np->cwd = idup(p->cwd);
safestrcpy(np->name, p->name, sizeof(p->name));
np->parent = p;

// Set state to RUNNABLE
np->state = RUNNABLE;

// return new thread's pid
return np->pid;
```

2.) join()

Starting the join() function with acquiring lock. This lock ensures that only one process at a time will traverse proc[]. Because there would be a race condition between join() and exit() so it is necessary to use lock.

```
// join:
int join (void **stack){
    struct proc *p;
    int children_, pid;

    // ensuring synchronized access to process table
    acquire(&wait_lock); // join() and exit() may cause race condition. | You, 1초 전 • Uncommitt
```

The join function uses a for loop that ends when a zombie thread is found, or there is no child thread, or the process is killed. And inside the loop, I create another loop that searches through proc[NPROC] array and filter the necessary processes.

children_ variable indicates number of child threads.

```

for(;;){
    children_=0; // will end the loop when zombie is found, no child thread, or current process is 1

    // process table loop
    for (p = proc; p < &proc[NPROC]; p++){ // linear search through proc[NPROC] array
        // skip unrelated processes
        if(p->parent != myproc() || p->pagetable != myproc()->pagetable)
            continue;

        acquire(&p->lock);
        children_ = 1;

```

If we find a zombie thread, and if then copy the value of `p->user_stack` into the user-space pointer that was passed to `join()` (so that in `thread_join()` we can call `free(stack)`), then release the lock and return failure signal. Otherwise, I free the trapframe and allocated struct `proc`, and return `pid`.

```

// looking for zombie thread, finished and called exit()
if (p->state == ZOMBIE){
    // copy thread's user stack pointer to caller
    *stack = p->user_stack;
    pid = p->pid;
    freeproc(p);
    release(&p->lock);
    release(&wait_lock);
    return pid;
}
release(&p->lock);
}

```

And this is to handle the case if there's no zombie thread. I set it to sleep until a child changes state to zombie.

```

// if there are children but no zombie -> sleep.
if(!children_ || killed(myproc)){
    release(&wait_lock);
    return -1;
}

sleep(myproc(), &wait_lock);
}
}

```

3.) `thread_create()` and `thread_join()`

I've added a new header and c file of thread in user space.

```

1  #ifndef THREAD_H
2  #define THREAD_H
3
4  int thread_create(void (*start_routine)(void *, void *), void *arg1, void *arg2);
5  int thread_join(void);
6
7  #endif
8

```

It calls the kernel join() system call, and then frees the memory for the stack that has received from join(). And then it returns the pid of the joined thread, or -1 if error case.

```

int thread_join()
{
    void *stack; // stack pointer to be freed

    int pid = join(&stack); // calls kernel join()

    if (pid < 0)
        return -1;

    // free the user stack after the thread has terminated
    free(stack);
    return pid;
}

```

Because the caller doesn't supply a stack so I've allocated a stack here. Stack is page-aligned because malloc() in xv6 returns page-aligned blocks by default. So if clone() fails then return -1, if it succeeds then returns the new thread's pid.

```

int thread_create(void (*start_routine)(void *, void *), void *arg1, void *arg2)
{
    void *stack = malloc(4096); // allocate 1 page for stack

    if (stack == 0)
        return -1;

    // Call the kernel's clone syscall
    int pid = clone(start_routine, arg1, arg2, stack);

    if (pid < 0) {
        free(stack); // cleanup if clone fails
        return -1;
    }

    return pid;
}

```

Modifications in Makefile:

```

ULIB = $U/ulib.o $U/usys.o $U/printf.o $U/umalloc.o $U/thread.o

```

4.) System Call Modifications

exec() replaces the current process's address space, but if this process is a thread then those sibling threads still exist and will now be running on a memory that's either freed proc_freepagetable() or replaced with a new layout. So the solution: before replacing the address space and assigning the new pagetable, find and kill all sibling threads except the calling thread. It ensures only the calling thread survives and enters the new program. Also, sibling threads get marked killed = -1 and exit as soon as scheduled.

```
for (struct proc *t = proc; t < &proc[NPROC]; t++) {
    if (t != p && t->pagetable == p->pagetable) {
        acquire(&t->lock);
        t->killed = 1;
        if (t->state == SLEEPING)
            t->state = RUNNABLE;
        release(&t->lock);
    }
}
```

The problem of sbrk() is because threads created by clone() share the same pagetable and heap size sz. myproc()->sz, the default usage might cause race condition. Thus, I modify growproc() to use global lock memlock for synchronization, and after resizing memory, update heap size sz for all threads sharing the same pagetable.

```
int
growproc(int n)
{
    uint64 sz;
    struct proc *p = myproc();

    acquire(&memlock);

    sz = p->sz;
    if (n > 0) {
        if ((sz = uvmmalloc(p->pagetable, sz, sz + n, PTE_W)) == 0) {
            release(&memlock);
            return -1;
        }
    } else if (n < 0) {
        sz = uvmdalloc(p->pagetable, sz, sz + n);
    }

    for (struct proc *t = proc; t < &proc[NPROC]; t++) {
        if (t->pagetable == p->pagetable) {
            t->sz = sz;
        }
    }

    release(&memlock);
    return 0;
}
```

I modify kill() in proc.c so that it kills all threads sharing the same pagetable.

```
int
kill(int pid)
{
    int found = 0;

    for (struct proc *p = proc; p < &proc[NPROC]; p++) {
        acquire(&p->lock);
        if (p->pid == pid) {
            // Kill all threads sharing the same address space
            pagetable_t target_pagetable = p->pagetable;
            release(&p->lock);

            for (struct proc *q = proc; q < &proc[NPROC]; q++) {
                acquire(&q->lock);
                if (q->pagetable == target_pagetable) {
                    q->killed = 1;
                    if (q->state == SLEEPING)
                        q->state = RUNNABLE;
                }
                release(&q->lock);
            }

            return 0;
        }
        release(&p->lock);
    }
    return -1;
}
```

III. Result & Troubleshooting

```
xv6 kernel is booting
init: starting sh
$ thread_test

[TEST#1]
Thread 0 start
Thread 1 start
Thread 1 end
Thread 0 join failed
$ Thread 2 start
Thread 2 end
Thread 3 start
Thread 3 end
Thread 4 start
Thread 4 end
Thread 0 end
□
```

I've failed from the first test, which I think is the problem occurring in system calls `clone()` or `join()`. Either of them (or probably both of them) are returning -1 before the thread becomes a zombie. So my `thread0` join is keep failing.

Right now I also am confused by the boundary. (the thread to block and not to block.) Thus it issued me a contiguous, non-stop of cloning.

Also the problems I met were locking problems. For a very long time, I've met "panic: acquire/release" issues. This is because I was confused by the timing of acquiring the lock, releasing the lock and the number of locks that I've used. One of the problems was that I didn't know `allocproc()` already contained acquiring the lock so I had to remove the line and manually acquire locks in my system calls `join()` and `clone()`.